

# Assignment: Build a RAG-Powered Chatbot for News Websites

## 1. Overview

This is an assignment for the role of Full Stack Developer at Voosh.

You are required to create a simple full-stack chatbot that answers queries over a news corpus using a Retrieval-Augmented Generation (RAG) pipeline. Every new user should be a new session so create a session identifier.

---

## 2. Objectives

### 1. RAG Pipeline

- Ingest ~50 news articles (e.g. RSS feed or scraped HTML).
- Embed with Jina Embeddings (free tier) or any other open source embeddings you like .
- Store embeddings in a vector store of your choice (Qdrant, Chroma, faiss, etc.).
- Retrieve top-k passages for each query, then call Gemini API for final answer.

### 2. Back-End

- Build a REST API (Node.js – Express).
- Use API / Socket based for chat. Also Implement to fetch session's history and clear session
- **Storage:**
  - **Redis** for in-memory chat history (per session).
  - Persist final transcripts optionally in a SQL database (MySQL/Postgres)(Optional).

### 3. Front-End

- React + SCSS with:
  1. **Chat screen:**
    - Displays past messages.
    - Input box for new messages.
    - Streaming bot responses (if possible) or typed-out reply.
    - button to reset the session

### 4. Caching & Performance

- Cache session history and conversations in In-memory database
  - Show in your README how you'd configure TTLs or cache warming.
-

### 3. Tech Stack (choose & justify)

- **Embeddings**
  - **Vector DB**
  - **LLM API:** Google Gemini (free trial)
  - **Backend:** Node.js (Express)
  - **Cache & Sessions:** Redis (in-memory) or any other in-memory database you like
  - **Database (optional):** MySQL or Postgres
  - **Frontend:** React + SCSS
- 

### 4. Deliverables

Please share your work at [richa@voosh.in](mailto:richa@voosh.in) with-

1. **List of Tech Stack Used**
  2. **Git Repositories**
    - Two public repos (frontend & backend) with full code and a clear `README.md` in each.
  3. **Demo Video**
    - A video (mp4 or hosted via unlisted link) showing:
      - Starting the frontend.
      - Sending queries and observing Gemini responses.
      - Viewing and resetting chat history of session.
  4. **Code Walkthrough**
    - Written or video explanation of the end-to-end flow, covering:
      - How embeddings are created, indexed, and stored.
      - How Redis caching & session history works.
      - How the frontend calls API/Socket and handles responses.
      - Any noteworthy design decisions and potential improvements.
  5. **Live Deployment**
    - A hosted, publicly accessible link (using any free hosting service) where we can test the chatbot.
-

## 5. Evaluation Criteria

| Area                    | Weight |
|-------------------------|--------|
| End-to-End Correctness  | 35%    |
| Code Quality            | 30%    |
| System Design & Caching | 20%    |
| Frontend UX & Demo      | 5%     |
| Hosting                 | 10%    |

**Good luck!**

We look forward to reviewing your working demo, code repos, and detailed flow explanation and testing it.

**Note:** Feel free to leverage any existing code—from open-source libraries, GitHub repositories, or AI-generated snippets—provided you fully understand and can explain it.

## Resources & References

- News ingestion example: <https://github.com/fhamborg/news-please>
- Reuters sitemaps: <https://www.reuters.com/arc/outboundfeeds/sitemap-index/?outputType=xml>
- Jina Embeddings: <https://jina.ai/embeddings>
- Google AI Studio API keys: <https://aistudio.google.com/apikey>
- Qdrant quickstart: <https://qdrant.tech/documentation/quickstart/>
- Pinecone quickstart: <https://docs.pinecone.io/guides/get-started/quickstart>
- [Render.com](https://render.com) hosting: <https://render.com/>
- Redis Python client: <https://github.com/redis/redis-py>
- For frontend development you can use - v0, [bolt.new](https://bolt.new) or any other LLM.